

Code

General Stuff

```
# limit number of digits:
options(digits=2)
# apply a function to vector
# compute corr of columns
cor.vec <- apply(x,2,cor,y=y)
# apply a function n times
replicate(n, cv.knn1(x.new,y,K=10))
# find the idx in order
order(c(3,6,2,8,4,9,1),decreasing=T)
#return bool vector not na rows
complete.cases(df)
#removes na rows from df
na.omit(df)
which.min # same as argmin
which # returns TRUE indices
rm(list=ls()) # clear everything
"X1" %in% c("X2", "X3") #TRUE if exists in it
# for making grids
seq(from,to,length,by)
# reshaping into 10by10 matrix
matrix(1:100, nrow=10)
```

```
# string manipulation
# paste is like paste0 but sep=" "
paste0(c("x","+", "y"),collapse="") # "x+y"
paste0("x","+", "y") # "x+y"
paste("x","+", "y") # "x + y"
paste("x","y", sep=" + ") # "x + y"
```

Linear Model

```
fit <- lm(y ~ x, data)
# 0 + x for no intercept
# functions
summary abline coef
residuals fitted confint
# get std of first feature
se1<-coef["x1", "Std. Error"]
plot(fit) # diagnostic plots
# conducting a partial F-test:
anova(fit,partialmodel, fit.all)
# prediction
df <- data.frame(colname = x0)
predict(fit, df, level=0.95,
interval=c("p", "c"))
# manual calculation
quant <- qt(.975,n-p)
sigma.hat <- sqrt(sum((fit$resid)^2)/(n-p))
X <- as.matrix(cbind(1,x))
XtXi <- solve(t(X) %*% X)
x_00 <- as.matrix(c(1,x_0),nrow=2)
# for confidence interval use
se<-sigma.hat*sqrt(t(x_00)%*%XtXi)%*%x_00
# for prediction interval use
se <-sigma.hat*sqrt(1+t(x_00)%*%XtXi)%*%x_00)
lower <- expected_y0 - quant*se
upper <- expected_y0 + quant*se
```

Matrix

```
t # transpose
%*% # matrix product
crossproduct
```

```
solve # inverse
as.matrix #to convert df to matrix
```

Plotting

```
par(mfrow=c(2,2)) # 2x2 grid
dev.off() # clear
pairs # pairs plot
plot(stuff, pch, xlim=c(-1,1),
xlab="X label",main="Main title")
legend("topleft", c(names), pch, col)
# pch/col is symbol/color for each entry
hist(breaks,freq=F,add=T)
abline(a, b, h, v , lwd, lty, col)
# plot distribution
lines(density(vector))
# Plot the columns of one matrix
# against the columns of another.
matlines(x, ys)
# for small data
stripchart(data,method="stack")
stem(data)
```

KNN

```
# Package for knn classification:
library(class)
# knn classifier
y.pred <- knn(x.train, x.test, y.train, k=1)
# Package for knn regression:
library(kknn)
fit <- kknn(Y ~ ., train_df, test_df, k=8)
# predictions on test data
pred <- predict(fit)
```

Cross Validation

```
#Automatic CV for linear models
library(boot)
lm.fit = glm(y ~ bs(x, df = i), data = df)
RSS[i] = cv.glm(df, lm.fit, K = 10)$delta[2]
# randomly shuffle the rows:
ind.x <- sample(1:n, replace=FALSE)
x <- x[ind.x,]
# create K folds:
folds <- cut(seq(1,n),breaks=K,
labels=FALSE)
ind.test <- which(folds==i)
```

Bootstrap

```
library(boot)
# non parametric bootstrap
mle.fn <- function(data, index, extra_par){
  return(estimator)
  #for example(regression)
  return(coef(lm(y~x,data=data,subset=index)))
}
boot.result <- boot(data, mle.fn, R,
extra_par=xx)$t
# empirical CDF
boot.scaled[,i] <-
sqrt(n)*(boot.result[,i] - mles[i])
ecdf(boot.scaled[,i])
# parametric bootstrap
# function to compute estimator
mle.fn <- function(data) {
  return(estimator) #like median
```

```
}
# function to generate data
data.gen <- function(data, mle) {
  out <- data
  out$hours <- rexp(nrow(out), mle)
  return(out)
}
# perform boot
air.boot <- boot(data, mle.fn, R=1000,
sim = "parametric",
ran.gen = data.gen,
mle = 1/mean(aircondit$hours))
boot.ci(boot.out,
type = c("norm", "basic", "perc"),
index=i)
# note normal CI uses bias correction.
# its midpoint is \hat \theta_n - \hat{bias}
#"norm" = normal, "basic" = reversed quantile
#"perc" = quantile, "stud" = bootstrap T
# Fixed X approach: bootstrap residuals
# Resampling here. Can be model-based.
boot.resid <- function(data, indices) {
  resid.new <- fit$residuals[indices]
  y.new <- fit$fitted + resid.new
  data.new <- data.frame(x=x, y=y.new)
  fit.new <- lm(y ~ x, data=data.new)
  return(coef(fit.new))
}
fit <- lm(y ~ x, data=data)
boot(data=data, statistic=boot.resid, R=1000)
```

Power Assesment

```
# compute p-value
gen.samples.adj <- c(gen.samples, value)
p.val <- sum(gen.samples.adj <=value) /
length(gen.samples.adj)
# compute rejection region
rej <- quantile(gen.samples, 0.05)-1
# compute power
power <- sum(gen.samples.alt <= rej) / nsimul
```

Permutations

```
combn(n,k) # choose k out of n
# e.g y1=control, y2=treatment
wilcox.test(y1, y2, alternative, paired)
# implicit paired signed rank test
wilcox.test(diff, alternative)
# signed rank test
V <- sum(rank(abs(data))[data>0])
```

Model Selection

```
library(leaps)
fit <- regsubsets(y~, df, nbest, nvmax,
force.in, force.out, intercept, method)
coef(fit, n) # best model with n predictors
fit.sm <- summary(fit)
# parameters used in each step:
fit.sm$which
# best model based on bic:
best <- which.min(fit.sm$bic)
params <- coef(fit, best)
# add intercept and handle categorical
```

```
test.mat <- model.matrix(y~, data=df[test,])
pred <- test.df[,names(params)] %*% params
plot(fit.sm, scale) # r2, adjr2, Cp, bic
```

```
# backward regression
# trace=1 for verbose
# k=log(n) gives BIC penalty
model <- stepAIC(y ~ ., data=df), k = log(n),
trace = 0)
```

Ridge and Lasso

```
library(glmnet)
grid <- 10^seq(from = 10, to = -2, length=100)
model <- glmnet(x, y, alpha=0, lambda=grid)
# alpha=0 -> ridge, alpha=1 -> lasso
plot(model, xvar="lambda")
# get coefficient estimates for new lambda=35
predict(model, s=35, type="coefficients")
# predict on test data
predict(model, s=35, newx=x.test)
#automatic cross validation
cv.out <- cv.glmnet(x, y, alpha=0)
plot(cv.out)
cv.out$lambda.min # best lambda
# +1 standard error larger lambda
cv.out$lambda.1se
# set specific folds(inner cv)
inner.folds <- folds[folds != i]
inner.folds <- factor(inner.folds)
levels(inner.folds) <- 1:(k-1)
cv.glmnet(...,
foldid = as.numeric(inner.folds),...)
```

NonLinearity

```
# poly returns uncorrelated columns
# raw=T -> without orthogonalization
fit <- lm(y~poly(x,4),data=df)
predict(fit,newdata=data.frame(x=grid))
```

```
library(splines)
# bs is basis function
fit <- lm(y~bs(x,knots=c(k1,k2,k3)),data=df)
# For K knots, specify df=K+3
fit <- lm(y~bs(x,df=K+3),data=df)
# locations of knots
attr(bs(age,df=6),"knots")
# bs() default degree is 3 (cubic splines)
# use ns() to fit natural splines.
# for K knots, specify df=K+1
fit <- lm(y~ns(x,df=4),data=df)
```

```
# Smooth splines
fit <- smooth.spline(x,y,df=16)
fit.cv <- smooth.spline(x,y,cv=TRUE)
fit.cv$df
fit.cv$lambda
```

```
# Local Regression
# span is how much nearby data to consider
fit <- loess(y~x,span=.2,data=df)
```

```
# GAMs
library(gam)
```

```
gam <- lm(y~ns(x1,4)+ns(x2,5)+x3,data=df)
plot.Gam(gam, se=TRUE, col="red")
```

```
# use backfitting
# s() for smoothing splines, lo() for loess
gam <- gam(y~s(x1,4)+s(x2,5)+x3,data=df)
```

```
# formula generation
# use all columns
list.var <- names(df)
f <- sapply(list.var,
function(x){paste0("s(",x,",4)")}
f <- paste0(f, collapse = "+")
f <- paste0("y ~ ", f)
f <- as.formula(f)
gam <- gam(f, data=df)
```

Trees

```
library(tree)
#.-y means everything but y
fit <- tree(y~.y, data=df, subset=train_idx)
plot(fit)
text(fit,pretty=0)
#prediction
tree.pred = predict(tree.fit, testdf,
type="class")
table(tree.pred, ytest)
cv.fit <- cv.tree(fit, FUN=prune.tree)
plot(cv.fit$size, cv.fit$dev,type="b")
# for classification:
# replace prune.tree to prune.misclass
# k is alpha
# dev means number of misclassifications
```

```
# best is number of terminals
pruned <- prune.tree(fit, best=2)
```

```
# Random Forests
library(randomForest)
fit <- randomForest(y~.,data, subset, mtry,
norm.votes,importance=T, maxnodes)
# fitlpredicted gives predicted values based on
#out of bag samples
```

```
# mtry number of vars checked on each split
# default is p/3 for regression
# default is sqrt(p) for classification
# given MSE is out-of-bag MSE
```

```
# importance
varImpPlot(fit)
randomForest::importance(fit)
```

```
# fraction of trees voted for each class:
fit$vote
# aggregate class probabilities
# and not predictions:
library(ranger)
rf <- ranger( as.factor(y) ~. , data=data,
mtry, probability=TRUE, min.node.size)
rf$predictions
```

Formulas

Linear Regression

$y = X\beta + \epsilon, \epsilon \sim \mathcal{N}(0, \sigma^2 \mathbb{I})$

$\hat{\beta} = (X^T X)^{-1} X^T y$

$P = X(X^T X)^{-1} X^T$

$\hat{\beta} \sim \mathcal{N}(\beta, \sigma^2 (X^T X)^{-1})$

$\hat{y} \sim \mathcal{N}(X\beta, \sigma^2 P)$

$e = y - \hat{y} \sim \mathcal{N}(0, \sigma^2 (\mathbb{I} - P))$

$\hat{\sigma}^2 = \frac{1}{n-p} \sum_{i=1}^n (y_i - x_i^T \hat{\beta})^2$

$\hat{\sigma}^2 \sim \frac{\sigma^2}{n-p} \chi^2_{n-p}$

Adjusted R score

$TSS = \sum_{i=1}^n (y_i - y_{mean})^2, RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$

$\frac{1}{n-1} TSS$ = sample variance of y_1, \dots, y_n

$\frac{1}{n-1} RSS$ = sample variance of e_1, \dots, e_n

$R^2 = 1 - \frac{RSS}{TSS}$

$R^2 = 0$ means model do nothing

$R^2 = 1$ model completely fits

$R^2_{adj} = 1 - \frac{RSS/(n-p)}{TSS/(n-1)}$

$F = \frac{(RSS_1 - RSS_2)/(p_2 - p_1)}{RSS_2/(n - p_2)}$

Std Errors and P values of beta

$Se(\hat{\beta}_k) = \sqrt{\hat{\sigma}^2 [(X^T X)^{-1}]_{kk}}$

$\frac{\hat{\beta}_j - 0}{\sqrt{\hat{\sigma}^2 [(X^T X)^{-1}]_{kk}}} \sim t_{n-p}$ under null

Type 1 and 2 errors

		True State of Nature	
		H_0 true	H_0 false
Decision	Accept H_0	Correct Probability = $(1 - \alpha)$	Wrong (type II error) Probability = β
	Reject H_0	Wrong (type I error) Level of significance = α	Correct Power = $1 - \beta$

Confidence Intervals

$(1 - \alpha)100\%$ confidence interval for $\mathbb{E}[y_0] = x_0^T \beta$:

$x_0^T \hat{\beta} \pm \hat{\sigma} \sqrt{x_0^T (X^T X)^{-1} x_0} \quad t_{1-\frac{\alpha}{2}, n-p}$

$(1 - \alpha)100\%$ confidence interval for y_0 :

$x_0^T \hat{\beta} \pm \hat{\sigma} \sqrt{1 + x_0^T (X^T X)^{-1} x_0} \quad t_{1-\frac{\alpha}{2}, n-p}$

Bias Variance Decomposition

$$\mathbb{E}[(y_0 - \hat{f}(x_0))^2] = \overbrace{(f(x_0) - \mathbb{E}[\hat{f}(x_0)])^2}^{\text{Bias}^2} + \overbrace{\mathbb{E}[(\hat{f}(x_0) - \mathbb{E}[\hat{f}(x_0)])^2]}^{\text{Variance}} + \overbrace{\sigma^2}^{\text{Noise}}$$

Bootstrap

Consistency with convergence rate α_n for $\hat{\theta}_n$:

$$\mathbb{P}[\alpha_n(\hat{\theta}_n - \theta)] - \mathbb{P}^*[\alpha_n(\hat{\theta}_n^* - \hat{\theta}_n)] \xrightarrow{p} 0 \text{ as } n \rightarrow \infty$$

not random-want to know random-approx via bootstrap

Consistency holds if $\alpha_n(\hat{\theta}_n - \theta) \sim \mathcal{N}(\theta, \sigma^2)$

Bias and Variance estimation for $\hat{\theta}_n$:

$bias(\hat{\theta}_n) \approx \mathbb{E}^*[\hat{\theta}_n^*] - \hat{\theta}_n \approx \hat{\check{\theta}}_n - \hat{\theta}_n$

$var(\hat{\theta}_n) \approx var(\hat{\theta}_n^*) \approx \frac{1}{B-1} \sum_{b=1}^B (\hat{\theta}_n^{*b} - \hat{\check{\theta}}_n)^2$

Bootstrap CIs:

- quantile(percentile): $[q_{\hat{\theta}_n^*}(\frac{\alpha}{2}), q_{\hat{\theta}_n^*}(1 - \frac{\alpha}{2})]$
- normal(normal): $\hat{\theta}_n \pm q_Z(1 - \frac{\alpha}{2}) \hat{sd}(\hat{\theta}_n)$
where $Z \sim \mathcal{N}(0, 1)$, $\hat{sd}(\hat{\theta}_n) = \sqrt{\hat{var}^*(\hat{\theta}_n^*)}$ **uses $\hat{\theta}_n - bias$ as midpoint**
- reversed quantile(basic): $[\hat{\theta}_n - q_{\hat{\theta}_n^* - \hat{\theta}_n}(1 - \frac{\alpha}{2}), \hat{\theta}_n - q_{\hat{\theta}_n^* - \hat{\theta}_n}(\frac{\alpha}{2})]$
- Bootstrap T(stud): $[\hat{\theta}_n - q_{\frac{\hat{\theta}_n^* - \hat{\theta}_n}{\hat{sd}(\hat{\theta}_n^*)}}(\frac{\alpha}{2}) \hat{sd}(\hat{\theta}_n), \hat{\theta}_n - q_{\frac{\hat{\theta}_n^* - \hat{\theta}_n}{\hat{sd}(\hat{\theta}_n^*)}}(\frac{\alpha}{2}) \hat{sd}(\hat{\theta}_n)]$

Wilcoxon signed rank test

permute signes and use this statistic:
 $V = \sum_i rank(|D_i|) \cdot \mathbb{1}_{D_i > 0}$

Multiple Testing

		True State of Nature		Total
		H_0 true	H_0 false	
Decision	H_0 is not rejected	U	T	m - R
	H_0 is rejected	V	S	R
	Total	m_0	$m - m_0$	m

- $FDP : Q = \frac{V}{R}, FDR := \mathbb{E}[Q]$
- $FWER := \mathbb{P}(V \geq 1)$ global null : $m = m_0$
- $FWER = FDR$ under the global null
- $FWER \geq FDR$ in general
- if $m_0 \geq 1$ and the tests are exactly at level $\alpha \Rightarrow FWER \geq \alpha$
- $FWER \leq \alpha m$
- FWER Control - Bonferroni:**
conduct m tests at level $\frac{\alpha}{m}$

- under global null, with all tests at level α , $FWER = 1 - (1 - \alpha)^m$ and with first order tylor: $FWER \approx m\alpha$
- in case of dependant tests(We want distribution of min(pvalues) under the **global null**): conduct **FWER Control-Westfall Young** and find $\sigma := \alpha$ quantile of $min p_j$ with permutation test under the global null and conduct all tests at level σ . For a dataset with n points and m variables: shuffle y column, compute p-value for each m variables (e.g Wilcoxon) store $min(p_1, ..., p_m)$ repeat.
- FDR Control - Benjamini Hochberg:** want to have FDR rate q? order p-values and find the largest i such that $p_{(i)} \leq \frac{i}{m} q$ reject null hyp. for all tests $j, j \leq i$ (**Only works for independant p-values**)

Beyond Linearity

- Orthogonal Basis:** when adding new orthogonal basis to a function: estimates for other coef will be the same, p-values will change. Anova will use the p-values of the largest model.
- Number of free parameters in **cubic spline** with k knots? $k + 1$ intervals with 4 parameters $\Rightarrow 4(k + 1)$ parameters. k knots with 3 constraints per knot: $3(k)$ to be reduced. So in total $4(k + 1) - 3k = k + 4$ parameters. (In R reduce one to omit intercept $k + 3$)
- Number of free parameters in **natural cubic spline** with k knots(same as cubic spline but linear outside of knots)? $k + 4 - 2 * 2 = k$ (In R use $k + 1$)
- Good Basis function for cubic spline?
 $\beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3 + \gamma_1(x - c_1)_+^3 + \gamma_2(x - c_2)_+^3 + ... + \gamma_k(x - c_k)_+^3$
- Smoothing Spline**
 $\hat{g} = argmin_{g \in G} \{ \sum (y_i - g(x_i))^2 + \lambda \int_a^b g''(x)^2 dx \}$ if $\lambda = 0$ interpolate all points(**n degrees of freedom**), if $\lambda = \infty$ LS solution(**2 degrees of freedom**). In general equivalent to **shrunkn version of natural cubic spline**.

Backfitting Algorithm repeat

for each predictor j do
 $\hat{g}_j \leftarrow S_j(y - \hat{\mu} - \sum_{k \neq j} \hat{g}_k)$
 $\hat{g}_j \leftarrow \hat{g}_j - mean(\hat{g}_j)$ where $mean(\hat{g}_j) = \frac{1}{n} \sum_i \hat{g}_j(X_{ij})$
end for
until Convergence